

# Proof of Stake



# The Role of PoW

- Bitcoin, Ethereum and similar systems are open, permissionless networks
- Anyone can participate
- The system must agree on some “canonical order” of transactions
- Think of this mechanism as being like some kind of “voting”

# Sybil attacks



# The Role of PoW

- Economic defense against sybil attacks
- Idea: participants prove ownership of \$X worth of computer hardware by running computations on it 24/7 and uploading results to the network
- Computation must be **hard to perform, easy to verify**

# PoS

- Instead of using computation as the limiting economic resource, use digital assets inside the system
  - I.e. bitcoin, ether, or whatever other coin inside the platform
- Proving ownership of digital assets is easy: just sign a digital signature

# Chain-based PoS

- Attempts to closely replicate PoW
- Usual algorithm: every coin holder has some chance per second of being assigned the right to create a block
  - Probability is proportional to coin balance



# RNG Manipulation

- PoS block maker selection can only be pseudorandom, based on data within the chain?
- Can the RNG be manipulated?
- Idea: predict who will make future blocks if you do A vs B. Choose the option that will maximize your future revenue, even if its immediate rewards are lower.
- Possibly serious centralization risk


# RNG Manipulation

- NXT:
  - $s(G) = \emptyset$
  - $s(B) = h(s(B.parent) ++ B.parent.proposer)$
- RANDAO
  - $s(B) = h(s(B.parent) ++ B.reveal)$
  - Reveal must be pre-hash-committed-to
- Low-influence functions (eg. N-block majority rule)




# Nothing at stake flaw


Vote on neither  
 $EV = 0$




Vote on A  
 $EV = 0.9$



Vote on B  
 $EV = 0.1$




Vote on both  
 $EV = 0.1 + 0.9 = 1$




# Penalty-based PoS


Vote on neither  
 $EV = 0$




Vote on A  
 $EV = 0.9$



Vote on B  
 $EV = 0.1$



Vote on both  
 $EV = 0.1 + 0.9 - 5 = -4$





# Casper overview

- **Deposit + penalty**-based proof of stake
- Anyone can join as a **validator** by submitting a **deposit** (in ETH)
- Validators are penalized for contradicting themselves
- Heavily based on traditional BFT theory (Lamport, DLS, PBFT, ByzPaxos, etc) but remodeled to be more “chain-based”

# Security claims

- The algorithm defines a notion of “finality”; if a block is finalized then it is set in stone in history and we can never go back
- **Accountable safety:** if two conflicting blocks get finalized, then at least  $\frac{1}{3}$  of validators may lose their entire deposits
- **Plausible liveness:** it's always possible to finalize more blocks with  $\frac{2}{3}$  of validators

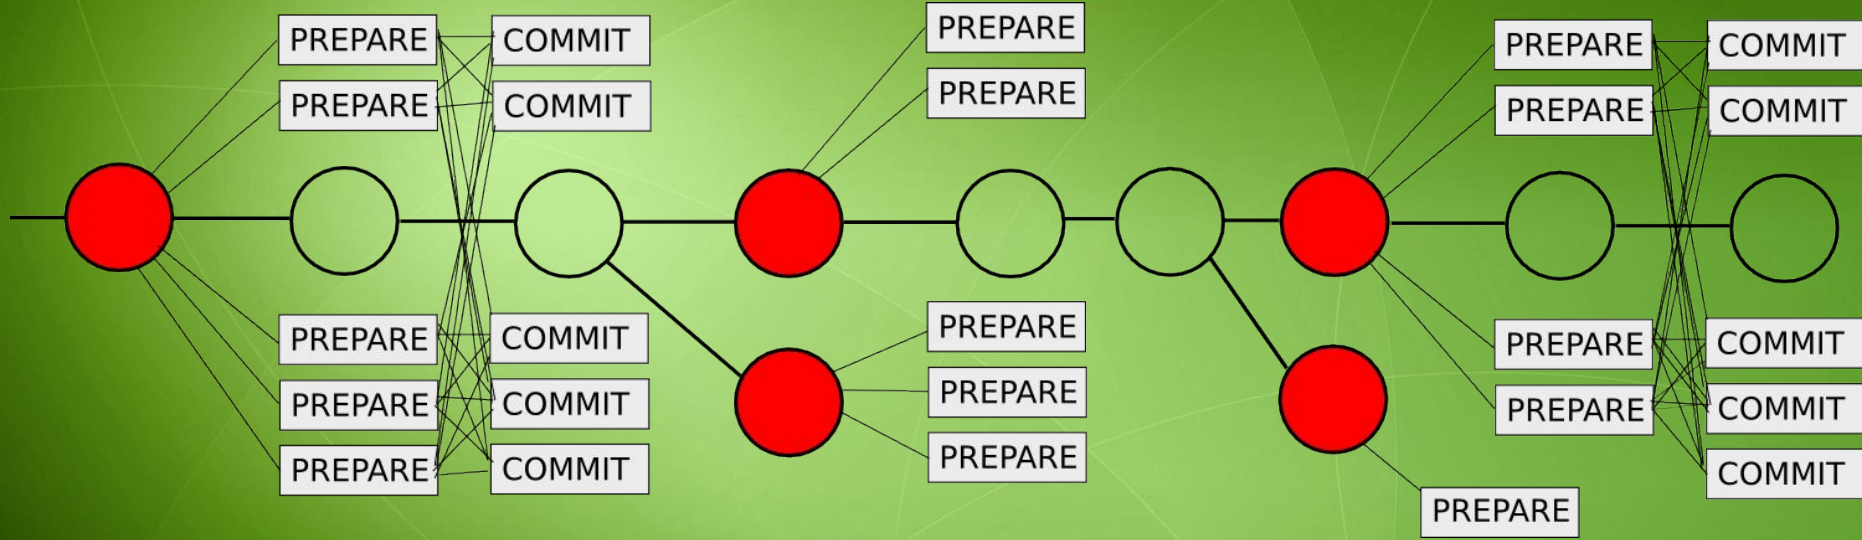


# Stage 1: Hybrid design

- Miners continue to create **proof of work blocks**
- Every 100th block is a **checkpoint**
- Active validators can send “prepare” and “commit” messages on checkpoints



# Stage 1: Hybrid design



# Justification and Finality

- The genesis is justified.
- If  $\frac{2}{3}$  of validators prepare a checkpoint C, all claiming some already justified checkpoint C' as a *source checkpoint*, then C is justified.
- If C is justified, and  $\frac{2}{3}$  of validators commit C, then C is finalized.

# Slashing conditions

- The protocol defines a set of **slashing conditions**, which represent actions that a validators is not supposed to do
- Example:
  - Sending a message when you are not supposed to
  - Sending two messages that contradict each other (eg. one votes for A, one votes for B)

# Slashing conditions

- If a validator violates a slashing condition, their deposit is deleted.

A slashing condition might look like this:

*If a validator sends a signed message of the form*

```
["PREPARE", epoch, HASH1, epoch_source1]
```

*and a signed message of the form*

```
["PREPARE", epoch, HASH2, epoch_source2]
```

*where `HASH1 != HASH2` or `epoch_source1 != epoch_source2`, but the `epoch` value is the same in both messages, then that validator's deposit is **slashed** (ie. deleted)*



```
# Cannot make two prepares in the same epoch
def double_prepare_slash(validator_index: num, prepare1: bytes <= 1000, prepare2: bytes <= 1000):
    # Get hash for signature, and implicitly assert that it is an RLP list
    # consisting solely of RLP elements
    sighash1 = extract32(raw_call(self.sighasher, prepare1, gas=200000, outsize=32), 0)
    sighash2 = extract32(raw_call(self.sighasher, prepare2, gas=200000, outsize=32), 0)
    # Extract parameters
    values1 = RLPList(prepare1, [num, bytes32, bytes32, num, bytes32, bytes])
    values2 = RLPList(prepare2, [num, bytes32, bytes32, num, bytes32, bytes])
    epoch1 = values1[0]
    sig1 = values1[5]
    epoch2 = values2[0]
    sig2 = values2[5]
    # Check the signatures
    assert ecrecover(sighash1,
                    as_num256(extract32(sig1, 0)),
                    as_num256(extract32(sig1, 32)),
                    as_num256(extract32(sig1, 64))) == self.validators[validator_index].addr
    assert ecrecover(sighash2,
                    as_num256(extract32(sig2, 0)),
                    as_num256(extract32(sig2, 32)),
                    as_num256(extract32(sig2, 64))) == self.validators[validator_index].addr
```



# Two slashing conditions

- NO\_DBL\_PREPARE:

CANNOT prepare two conflicting C1, C2 in one epoch

- COMMIT\_PREPARE\_CONSISTENCY

If you prepare in epoch  $e_p$ , with source in epoch  $e_s$   
then you cannot commit in any epoch  $e_c$  where  
 $e_s < e_c < e_p$



Yoichi Hirai

Follow

a convenience logician

Feb 25 · 7 min read

## Formal methods on some PoS stuff

In Paris, I received a description of a distributed consensus mechanism. If the description below looks ambiguous or impenetrable, this post is for you. I banged my head to formal verification tools called Alloy and Isabelle/HOL to clarify my understanding (code).

Message types:

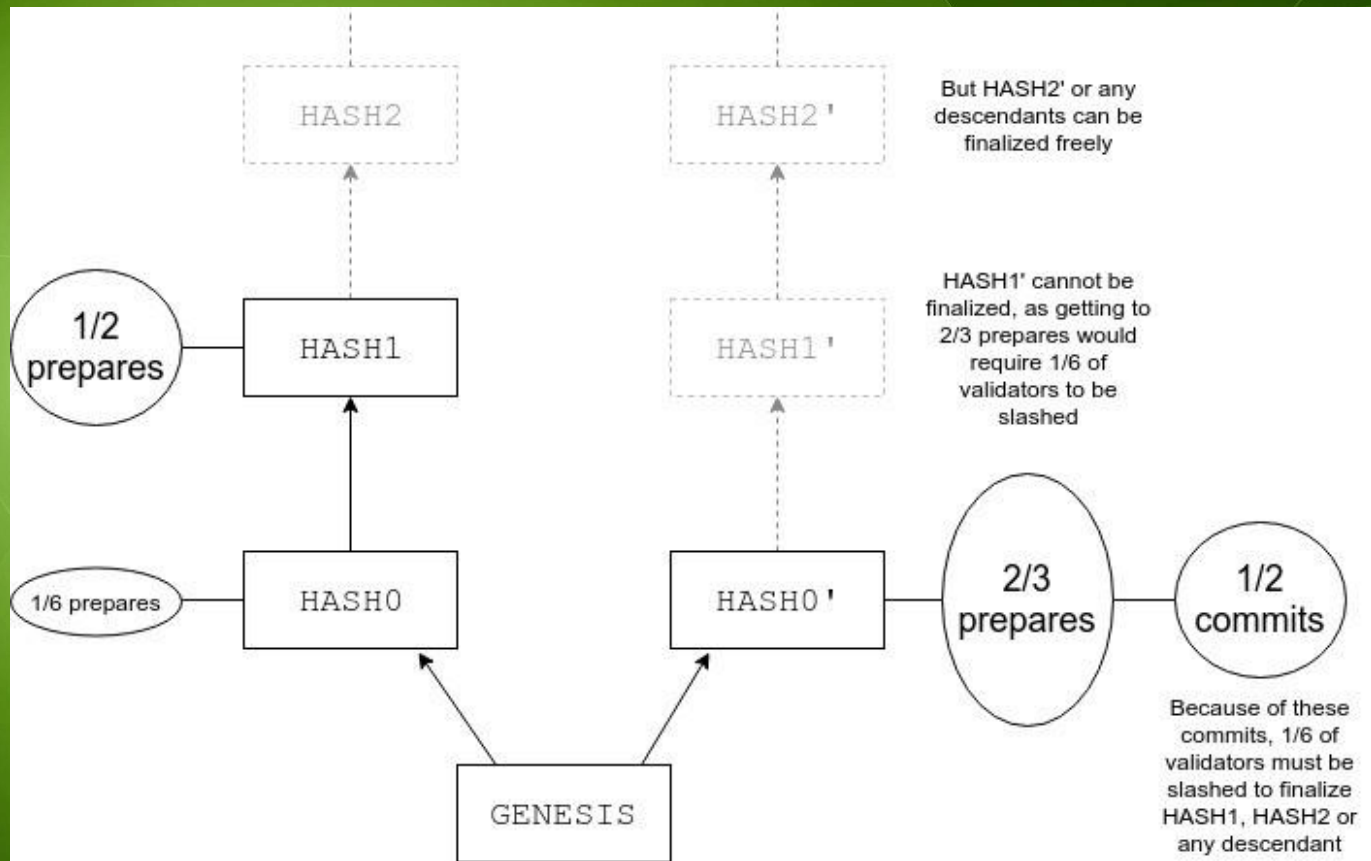
```
* commit(HASH, view), 0 <= view
* prepare(HASH, view, view_source), -1 <= view_source < view
```

Slashing conditions:

```
* [i] commit(H, v) REQUIRES 2/3 prepare(H, v, vs) for some
consistent vs
```

# Fork choice rule

- PoW: longest chain
- For hybrid PoS, PoW miners following the longest chain does not work



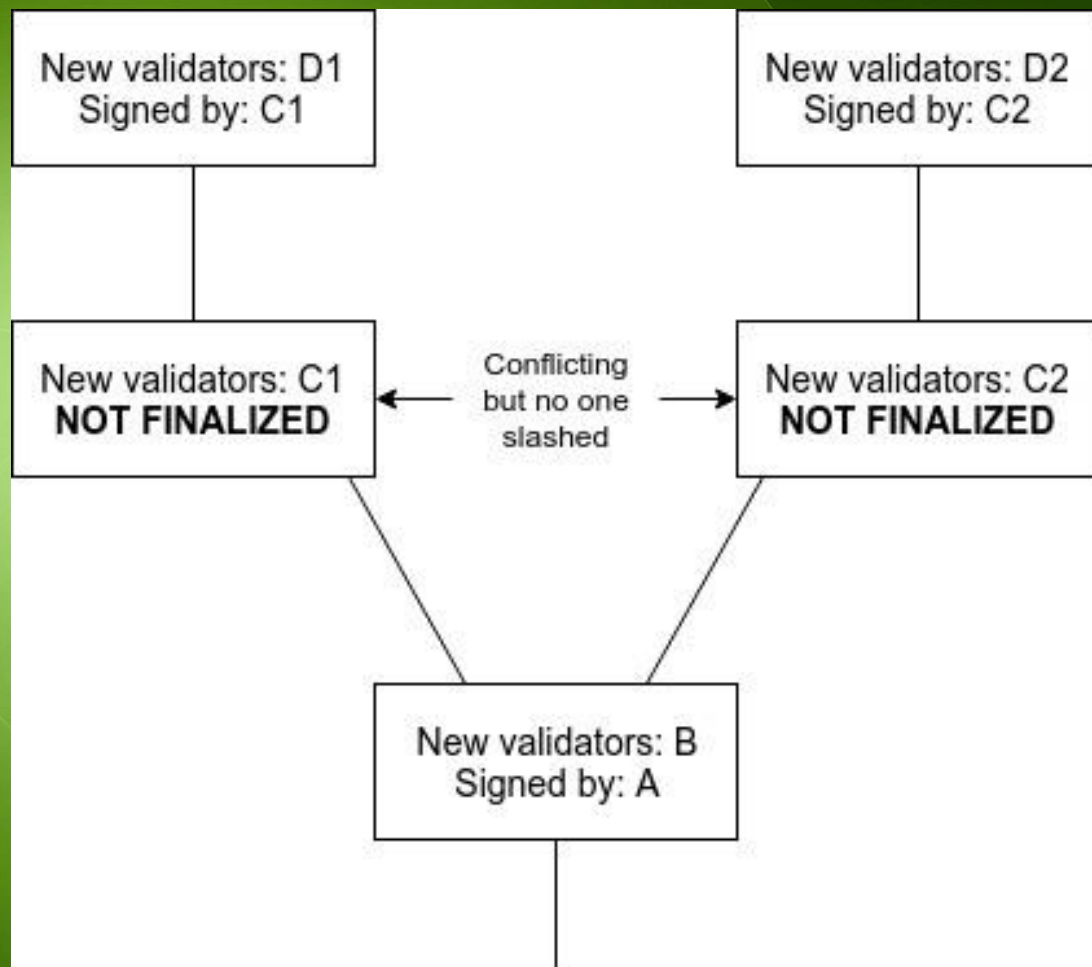
# Fork choice rule

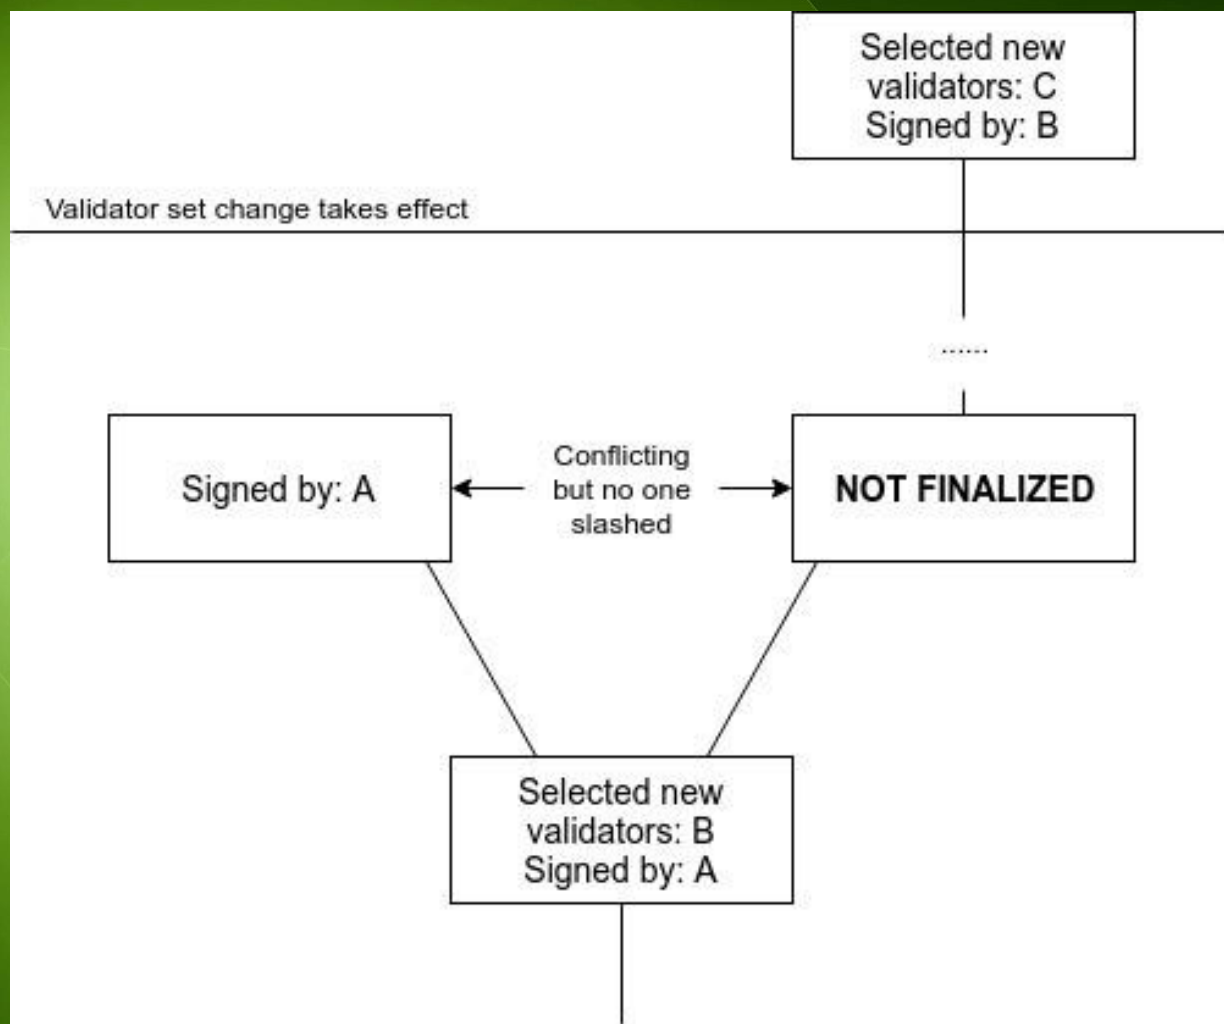
- For hybrid PoS, we can create a **hybrid fork choice rule**, roughly:
  - a. Prefer finalized checkpoints
  - b. Prefer checkpoints that are close to being finalized (ie. have  $\frac{2}{3}$  prepares and some commits)
  - c. Otherwise, take the PoW longest chain

# Dynamic validator sets

- How to deal with the possibility of the validator set changing?









Yoichi Hirai

Follow

a convenience logician

Mar 18 · 4 min read

## A mechanized safety proof for PoS with dynamic validators

I used a theorem prover Isabelle/HOL to check Vitalik's post about a proof-of-stake protocol that uses dynamic validator sets. (If you haven't seen, I did something similar for the simpler proof-of-stake protocol where the validator set is constant.) The proof script is available online.

# Dynamic validator sets

- If a validator leaves, they need to wait 4 months before they can recover their deposit
- Waiting period ensures that validators can be held accountable for slashing condition violations

# Implementation

- Casper contract (  
[https://github.com/ethereum/casper/blob/master/casper/contracts/simple\\_casper.v.py](https://github.com/ethereum/casper/blob/master/casper/contracts/simple_casper.v.py) )
- Casper fork choice rule (work being done on pyethereum)
- Daemon for sending prepares/commits (work being done in python; should be client-agnostic)



# Cryptoeconomic analysis



(applies to any PoS)



# Attacks

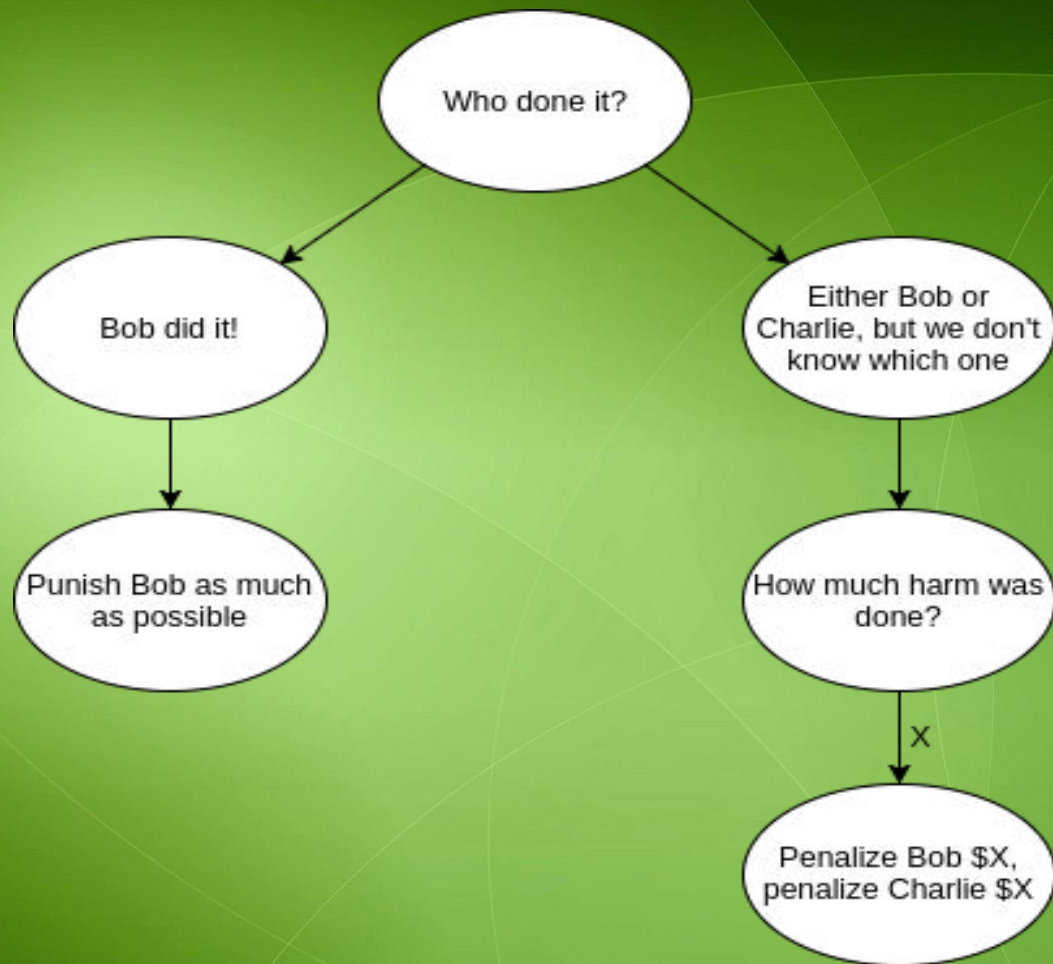
- Finality reversion attack (33%+)
- Going offline (33%+)
- Censorship (51%+)
- Extortion attacks (51%+)
- Discouragement attacks (any %)

# Fault categorization

- Outbound message omission
  - “I’m not going to vote”
- Inbound message omission
  - “Vote? What vote? I didn’t hear about a vote”
- Equivocation
  - “I vote A!” “I also vote B!”
- Invalid message faults
  - “I vote Harambe!”

# Fault categorization

- Uniquely attributable faults
- Non-uniquely attributable faults



Liveness is a  
non-uniquely-attributable fault





# Incentivization vs Niceness


- Penalize attackers: good
- Penalize innocent validators: bad
  - (because otherwise no one will want to validate due to risk)
- This is a tradeoff

# Subjective resolution

- If a serious attack happens, we can always resolve it subjectively (UASF, UAHF, etc)
  - Treating the community as a magic oracle that tells us who was the attacker
- But we want to use this as little as possible, as it harms decentralization

# Dealing with censorship





The background is a solid green color with a subtle gradient. Overlaid on this are several thin, light green lines that form large, overlapping circles or arcs, creating a geometric pattern.

Future directions



# Future directions

- Continuing to improve incentive structures
- Find a mechanism that combines together benefits of chain-based and BFT-based PoS
- Scalability (sharding etc)
- Speed up block times